

LeCoach

Your Private AI Audience

Local presentation rehearsal
with practical, supported feedback

Workplace AI
Battlefield Lightning · Global Online Track



Private rehearsal for workplace presenters

A workplace update, client pitch or internal demo

The presenter knows the material.
The harder rehearsal question is how
the delivery might affect an audience.

PRACTICING ALONE

Few cues about rushing
or staying turned toward
notes.

THE INTENDED BENEFIT


A specific passage to retry
and a concrete habit to
change.

A visible audience during rehearsal

The audience responds to sustained delivery signals.


Its timeline makes changes easy to revisit after a rehearsal.

Reactions simulate a practice audience. They do not measure real feelings.



Synthetic input · Computed audience · Devices off

The rehearsal loop



A completed session links a timestamp to an observation and an action.

Short or incomplete sessions show limitations when support is missing.

Local speech and approximate camera signals

Speech

Local Whisper transcription
Pace from finalized English words
Pauses and configured filler terms

Pace and filler estimates are approximate.
Filler recall still needs investigation.

Camera

Local MediaPipe pose processing
Approximate direction toward camera
Visible movement over time

Direction does not measure eye contact.
Unavailable observations stay unknown.

Feedback names a moment and a next action


Generated coaching from synthetic observations

Capture time	Supported observation	Practical rehearsal action
00:10 Improve	Pace about 200–205 WPM	Repeat more slowly. Pause after key points.
00:25 Improve	Sustained facing estimate of 0.2	Face the camera for the next key sentence.
00:40 Strength	144 WPM and cited facing estimate 0.85	Repeat the steadier, camera-facing delivery.

Every selected moment retains its supporting event IDs.
The selector reports fewer moments when evidence is limited.

One local pipeline

On the presenter's computer



The recorder also retains source observations.
Coaching explains the engine's evidence rather than creating new audience states.

Local processing with limited retention

**Raw audio and video recording
are off by default.**

Recognition, session events and feedback stay on the host.
The browser uses the loopback API.

The recorder keeps the active/completed session in memory
until the next session or application exit.

An actual vision-supported suggestion

00:04

Capture anchor

Sustained approximate
facing-away observations

Audience decision at 00:07.025
Neutral to bored

Next action

**Deliver your next key sentence
facing the camera, with your
notes beside it.**

Actual local CPU pickup · Vision only
No usable speech in this session

Sanitized run: see source references

Isolated camera timing on the Stage I host

Three fresh 5-second camera sessions on local CPU

Ubuntu 26.04.1 · Ryzen AI 7 450 · About 30 GiB RAM

Session	Startup (s)	Camera open (ms)	Inference (FPS)	Release (s)
1	0.3599	282.48	9.96	0.0774
2	0.2044	157.60	9.95	0.0708
3	0.2118	158.41	9.95	0.0541

Each session released the camera. No read or inference errors occurred.

This measures isolated vision behavior, not end-to-end UI latency.

Failures produce limitations and clean retries

Situation	Delivered behavior	Evidence basis
Missing microphone or camera	Show input notices and use supported observations only	Implementation + operator report
Too short or empty	Report insufficient evidence without invented moments	Computed synthetic checks
Stop and start again	Release inputs and keep sessions separate	Implementation + operator report
Late or incomplete input	Retain eligible drain observations and disclose gaps	Computed synthetic checks

Operator reports cover transcription, audience changes and feedback after stop. They do not establish quantified recognition accuracy or learning outcomes.

Expected value and pilot evaluation

A more focused next rehearsal

A presenter can choose one supported suggestion and retry the same workplace passage privately.

Proposed workplace pilot

Compare the first attempt with a retry.
Ask whether the cited moment and action help.
Track setup failures and unsupported advice.

A three-minute English demonstration

Planned video time	What the viewer sees
0:00–0:35	The rehearsal problem and the actual input mode
0:35–1:35	A rushed, notes-facing passage followed by a steadier retry
1:35–2:15	Actual audience timeline and supported feedback after stop
2:15–2:40	Local CPU processing and the Stage II hardware plan
2:40–2:55	One useful action for the next rehearsal

Live footage needs its own supported moments and visible mode label.
Final narration/edit timing and recording remain pending.

Stage I CPU delivery and Stage II hardware

Stage I

Local CPU prototype

Speech and camera processing
One engine, recorder and coach
English rehearsal experience

Stage II after qualification

UGen300 remains unmeasured

Verify USB/runtime/model compatibility.
Measure producer delays and UI delivery.
Test concurrent inputs and cleanup.

Hardware acceleration must preserve the same event contracts and evidence.

Source code and evidence

LeCoach

github.com/crasni/LeCoach

[Architecture and local retention](#)

[Synthetic scenario and generated coaching](#)

[Actual CPU functional report](#)

[Limited actual coaching pickup](#)

[Isolated CPU camera timing](#)

[Official ASUS UGen AI League competition rules](#)

Next validation: reviewed presenter/browser coaching and a timed English demo.